

AWS Cloud Practitioner Certification Bootcamp

Week - 5

Session 5 - Cloudformation, DevOps, CloudWatch, Cloudtrail

12th Feb, Saturday
7:00 PM to 8:30 PM IST



Speakers



Sanchit Jain

Lead Architect - AWS at Quantiphi
AWS APN Ambassador

Agenda



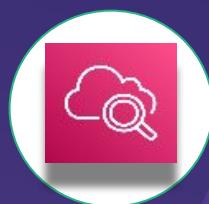
Re-cap of
Last session



AWS Cloudformation



DevOps on AWS



AWS CloudWatch



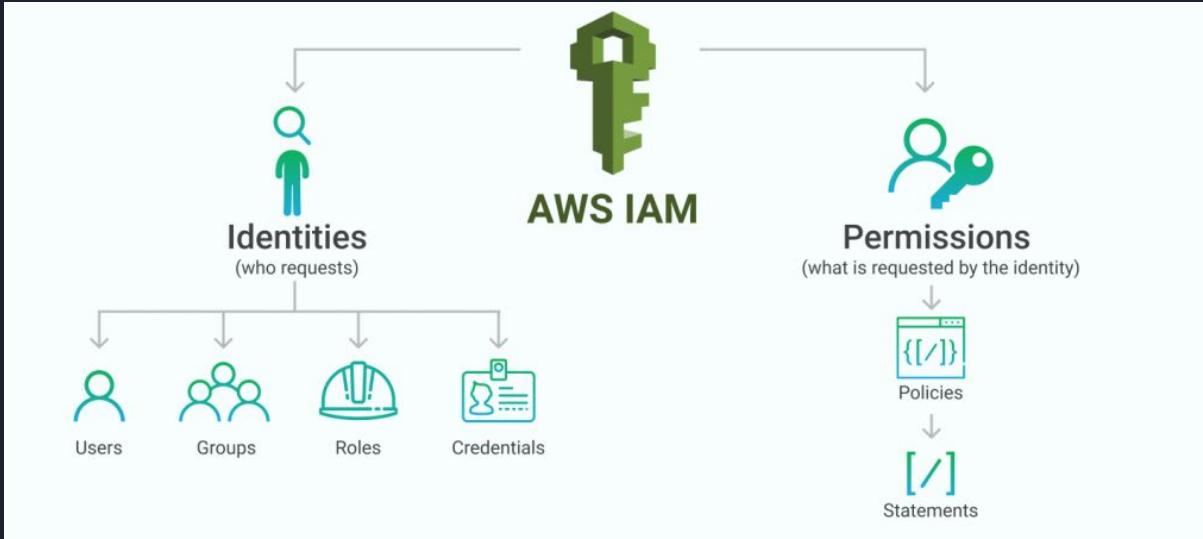
AWS CloudTrail



Re-cap of
Last session

AWS IAM - Users and Groups

- Root account created by default, shouldn't be used or shared as it has Administrative privileges.
- An **IAM user** is an entity that you create in AWS to represent the person or application that uses it to interact with AWS. A user in AWS consists of a name and credentials.
- Users performing same tasks with the same resources can be added to a Group.



IAM policy Evaluation Logic



AWS IAM Security Tools

IAM Access Advisor (user-level)

- Access advisor shows the service permissions granted to a user and when those services were last accessed.
- You can use this information to revise your policies.

The screenshot shows the AWS IAM Access Advisor interface. The top navigation bar includes 'AWS', 'Services', 'Edit', and a user profile 'zackAWS @ vc'. The left sidebar has links for 'Dashboard', 'Search IAM', 'Details', 'Groups', 'Users', 'Roles', 'Policies' (which is selected), 'Identity Providers', 'Account Settings', and 'Credential Report'. The main content area has tabs: 'Policy Document', 'Attached Entities', 'Policy Versions', and 'Access Advisor' (which is active). A note states: 'Access advisor shows the service permissions granted to this user and when those services were last accessed. You can use this information to revise your policies. This table does not include activity in the AWS São Paulo and Seoul regions. Learn more'. Below it says 'Note: recent activity usually appears within 4 hours. Access Advisor tracking began on Oct 1, 2015'. A search bar with 'Filter: No filter' and a 'Search' button is present. A table lists service names, access entities, and last accessed dates. The table header is 'Service Name' (sorted by click), 'Access by Entities' (highlighted with a red border), and 'Last Accessed'. The data is as follows:

Service Name	Access by Entities	Last Accessed
AWS Identity and Access Management	zackAWS and 1 more	Today
Amazon Elastic MapReduce	zackAWS	Today
Amazon CloudWatch	zackAWS	Today
AWS Security Token Service	zackAWS	Today

IAM Credentials Report (account-level)

- a report that lists all your account's users and the status of their various Credentials

AWS IAM Credential Report

1 Root Account
Ensure: no usage since last check, multi-factor authentication is enabled, and no access keys exist.

	A	B	C	D	E	F	G	H	I
1	user	user_creation_time	password_enabled	password_last_used	password_last_changed	mfa_active	access_key_1_active	access_key_1_last_rotated	access_key_1_last_used_date
2	<root_account>	2019-07-15T14:44:33	not_supported	2019-07-17T04:49:39	not_supported	TRUE	FALSE	N/A	N/A
3	pam_beasley	2019-11-13T18:32:34	FALSE	N/A	N/A	TRUE	TRUE	2020-06-18T12:12:27	2021-02-06T05:37:00
4	darryl_philbin	2021-01-25T19:12:26	FALSE	N/A	N/A	TRUE	TRUE	2021-11-25T16:11:26	N/A
5	dwight_schrute	2021-01-25T19:10:51	TRUE	2021-02-02T19:12:52	2021-01-25T19:10:51	TRUE	TRUE	2021-02-02T19:12:26	2021-02-02T03:31:00
6	kelly Kapoor	2021-01-25T19:13:23	TRUE	no_information	2021-01-25T19:13:23	FALSE	FALSE	N/A	N/A
7	ryan Howard	2021-01-25T19:26:22	TRUE	no_information	2021-01-25T19:26:22	FALSE	FALSE	N/A	N/A

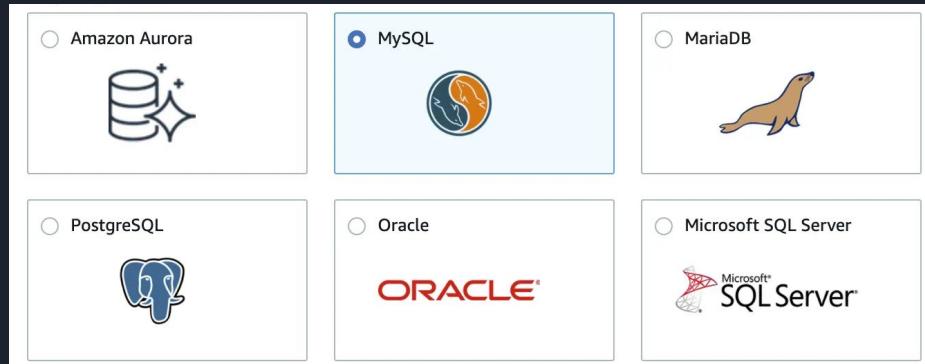
2 Users should have access keys or passwords, not both.
3 Delete user accounts that have never been used.
4 Password's should be changed based on company policy.
5 Enable MFA for all users.
6 Rotate old access keys.
7 Delete unused access keys.

IAM: Guidelines & Best Practices

- Lock away your AWS account root user access keys
- Create individual IAM users
- Assign users to groups and assign permissions to groups
- Grant least privilege
- Configure a strong password policy for your users
- Use and enforce the use of Multi Factor Authentication (MFA)
- Create and use Roles for giving permissions to AWS services
- Never share IAM users & Access Keys
- Rotate credentials regularly
- Remove unnecessary credentials
- Use policy conditions for extra security
- Audit permissions of your account with the IAM Credentials Report

Amazon RDS - Instances

- No special handling required. Just point to the new Instance
- A *DB instance class* determines the computation and memory capacity of a DB instance.
- Amazon RDS currently supports the following *DB engines*:



- Each DB engine has a set of parameters in a *DB parameter group* that control the behavior of the databases that it manages.

Amazon Aurora Databases

- Fully managed relational database engine compatible with **MySQL** and **PostgreSQL**. Includes high-performance distributed storage across AZs.
- Aurora maintains 6 copies of data across 3 AZs for Backup and Recovery.
- Has a master node and Read Replicas(upto 15). Replicas can Auto Scale based on Requests
- Aurora Endpoints
 - Writer endpoint - Points to master node
 - Reader endpoint - Common endpoint for all read-replica(s)
 - Custom endpoints - User defined, for offloading specific heavy analytical tasks

5X MySQL™

faster than standard RDS instances

3X 

Amazon Aurora Serverless

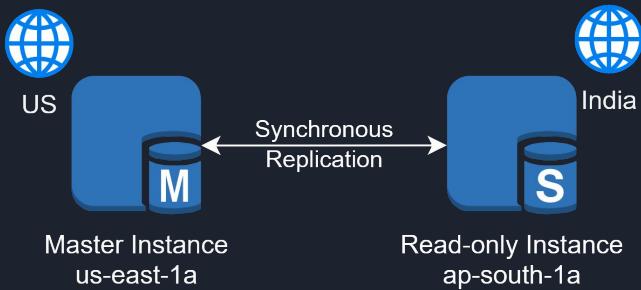
- Configuring capacity correctly in advance isn't always possible with the provisioned model. It can also result in higher costs if you overprovision and have capacity that you don't use.
- Checkout Serverless - Create a database endpoint without specifying the DB instance class or size !?
- Auto-Multi-AZ Failover, Always encrypted database



INFINITE POWER!

Amazon Aurora Global

- Enables low latency global reads and fast recovery from the rare outage that affect an entire AWS Region.



- The secondary cluster is read-only, so it can support up to 16 read-only Aurora Replica instances. Note you cannot auto-scale the read-replicas here

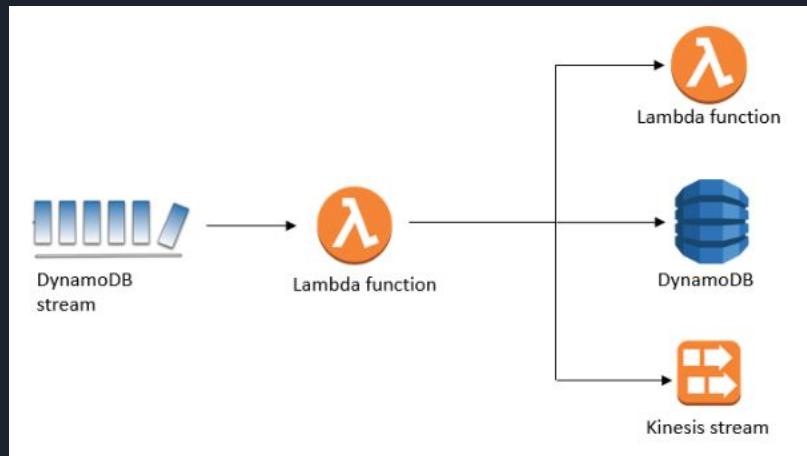
DynamoDB - Introduction

- Fully managed, fast NoSQL key-value database service
- No hardware provisioning, setup and configuration, replication, software patching, or cluster scaling required
- Allows to delete expired items from tables automatically to help reduce storage usage using TTL



DynamoDB Table Structure

Image Ref: AWS Blogs

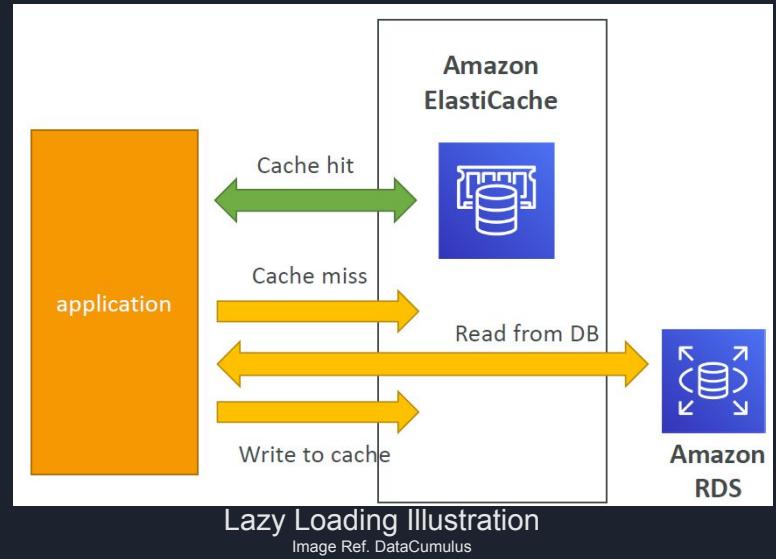
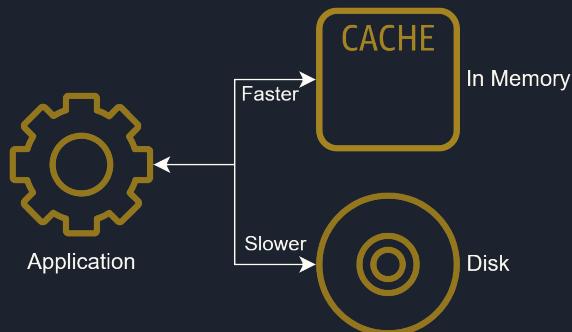


Streams Illustration

Image Ref: AWS Blogs

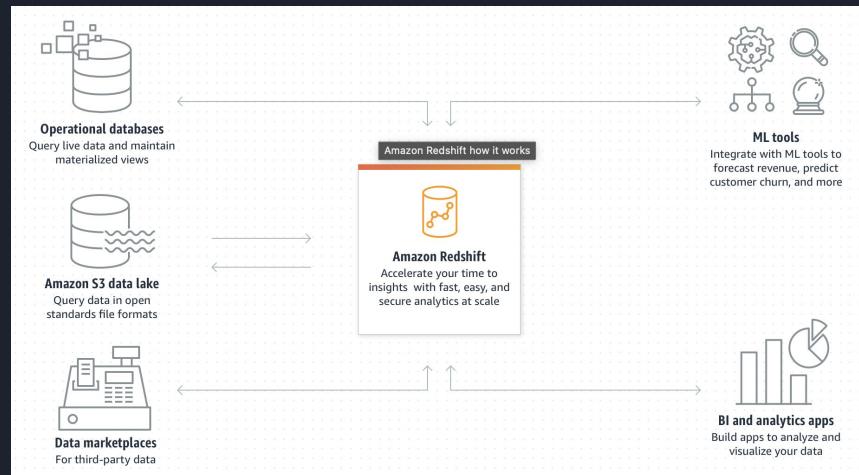
Elasticache - Introduction

- Provides a high-performance, scalable, and cost-effective caching solution and removes the complexity associated with deploying & managing a distributed cache environment.



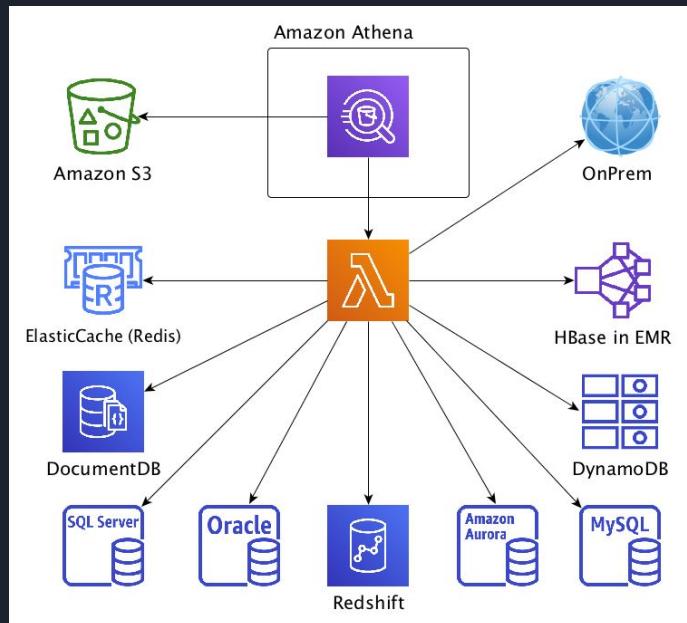
Redshift - Introduction

- Fully managed, petabyte-scale data warehouse service.
- An Amazon Redshift cluster is a set of nodes, which consists of a leader node and one or more compute nodes.
- Based on PostgreSQL but not used for OLTP workloads but rather used for OLAP (warehouse).
- Columnar storage, Massively Parallel Query Execution(MPP)
- COPY command helps to load data from S3 and similarly UNLOAD helps to offload data into S3.



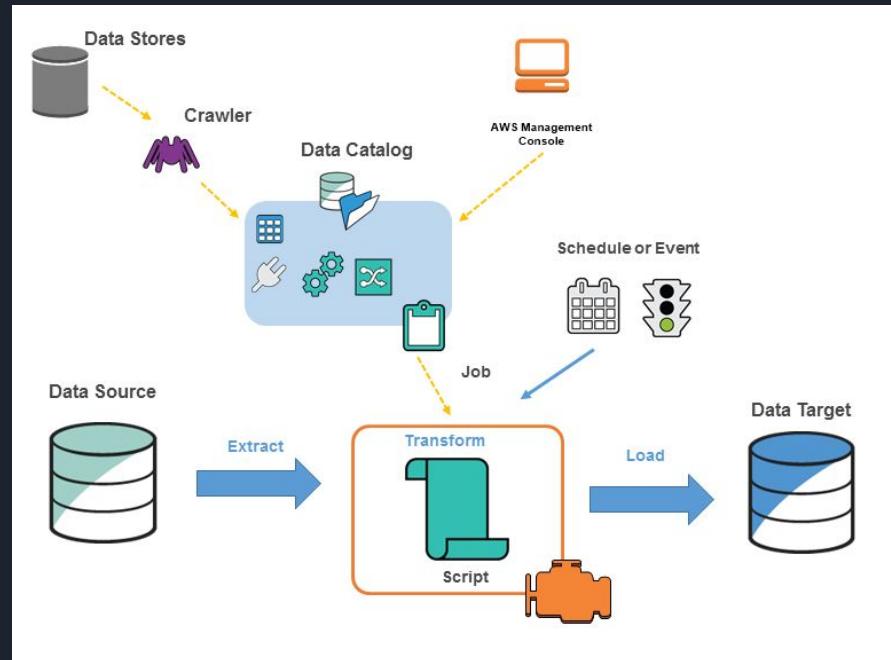
Amazon Athena

- Serverless query service to perform analytics against S3 objects
- Uses standard SQL language to query the files
- Supports CSV,JSON,ORC,Avro, and Parquet(builtonPresto)
- Pricing: \$5.00 per TB of data scanned
- Use compressed or columnar data for cost-savings (less scan)
- Use cases: Business intelligence / analytics, analyze & query VPC Flow Logs, ELB Logs, CloudTrail trails, etc...
- Exam Tip: analyze data in S3 using serverless SQL, use Athena



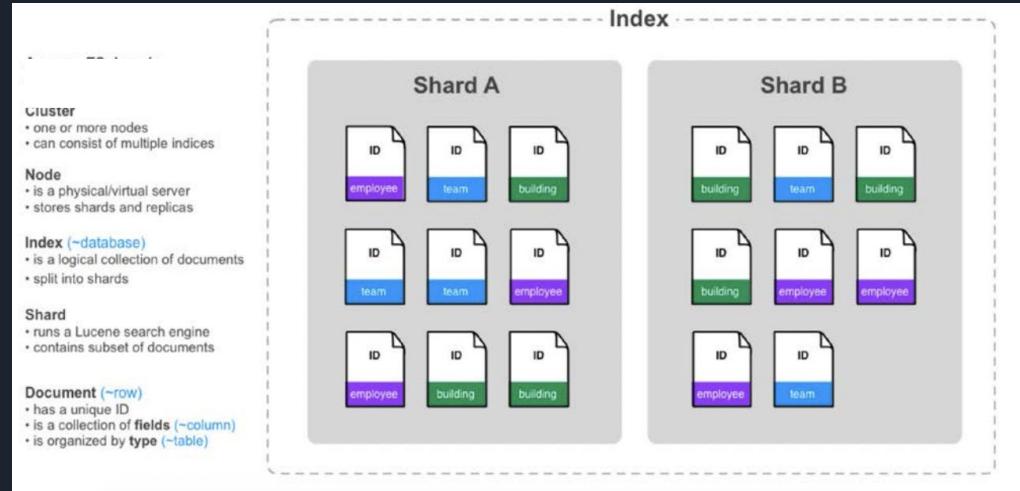
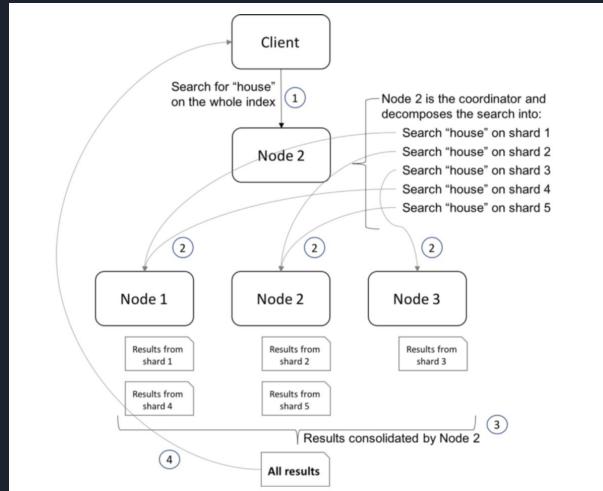
AWS Glue

- Fully managed serverless ETL service
- Serverless; runs on a fully managed, auto-scaling Spark environment.
- Glue Components
 - Databases - Tables and Connections
 - Data Catalog - Crawler and Classifiers
 - ETL - Jobs and Triggers



OpenSearch

- Fully open-source search and analytics engine for use cases such as log analytics, real-time application monitoring, and clickstream analysis
- Data visualization using OpenSearch Dashboards (the successor to Kibana)
- Also works on pattern matching(like)/partial matches of data quite fast.



AWS Cloudformation

CloudFormation Overview



All the resources required by user in an application can be deployed easily using templates

Also, you can reuse the template to replicate your infrastructure in multiple environments

To make the templates reusable, use the parameters, mappings, and conditions sections, in the template so that you can customize your stacks when you create them

CloudFormation Overview

- Infrastructure as a code(IAAC)
- One Click Creation and Deletion
- Declarative in nature
- Code can be version control
- Supports Yaml/Json



Image Credits:

[https://memegenerator.net/instance/72724467
/why-not-both-json-or-yaml-we-love-cloudformation](https://memegenerator.net/instance/72724467/why-not-both-json-or-yaml-we-love-cloudformation)

Template Structure

```
1 AWSTemplateFormatVersion: 2010-09-09
2 Description: >-
3   Launch an EC2 Instance running apache and expose
4   default page to the internet. Hardcoded to work
5   only with US-EAST-1 (N. Virginia)
6 Parameters:
7   InstanceType:
8     Description: WebServer EC2 instance type
9     Type: String
10    Default: t2.micro
11    AllowedValues:
12      - t2.nano
13      - t2.micro
14 Resources:
15   WebServer:
16     Type: 'AWS::EC2::Instance'
17     Properties:
18       Tags:
19         -
20           Key: Name
21           Value: Apache Default WebServer
22     InstanceType: !Ref InstanceType
23     # You shouldn't hardcode, just show here for example
24     ImageId: 'ami-0b898848883850657'
25     SecurityGroupIds:
26       - !GetAtt SecurityGroup.GroupId
27     UserData:
28       'Fn::Base64':
29         !Sub
30           #!/usr/bin/env bash
31           su ec2-user
32           sudo yum install httpd -y
33           sudo service httpd start
34     SecurityGroup:
35       Type: 'AWS::EC2::SecurityGroup'
36       Properties:
37         GroupDescription: Enable internet users to access.
38         SecurityGroupIngress:
39           - IpProtocol: tcp
40             FromPort: 80
41             ToPort: 80
42             CidrIp: 0.0.0.0/0
43 Outputs:
44   PublicIp:
45     Value: !GetAtt WebServer.PublicIp
```

Template Sections

- MetaData** Additional information about the template
- Description** A description of what this template is suppose to do
- Parameters** Values to pass to your template at runtime
- Mappings** A lookup table. Maps keys to values so you change your values to something else
- Conditions** Whether resources are created or properties are assigned
- Transform** Applies macros (like applying a mod which change the anatomy to be custom)
- Resources*** A resource you want to create eg. IAM Role, EC2 Instance, Lambda, RDS
- Outputs** Values that returned eg. an ip-address of new server created.

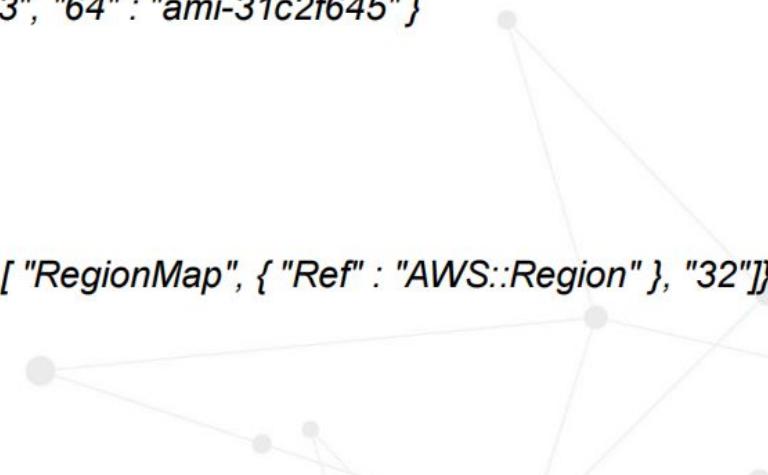
CloudFormation Templates requires you to **at least list one resource**.

Conditions

```
{  
    "AWSTemplateFormatVersion": "2010-09-09",  
    "Parameters": {  
        "EnvType": {  
            "Description": "Environment type.",  
            "Default": "test",  
            "Type": "String",  
            "AllowedValues": ["prod", "test"],  
            "ConstraintDescription": "must specify prod or test."  
        },  
        "Conditions": {  
            "CreateProdResources": {  
                "Fn::Equals": [  
                    {  
                        "Ref": "EnvType"  
                    },  
                    "prod"  
                ]  
            },  
            "Resources": {  
                "EC2Instance": {  
                    "Type": "AWS::EC2::Instance",  
                    "Properties": {  
                        "ImageId": "ami-0ff8a91507f77f867"  
                    }  
                }  
            }  
        }  
    }  
}
```

Mappings

```
{  
...  
"Mappings" : {  
    "RegionMap" : {  
        "us-east-1" : { "32" : "ami-6411e20d", "64" : "ami-7a11e213" },  
        "us-west-1" : { "32" : "ami-c9c7978c", "64" : "ami-cfc7978a" },  
        "eu-west-1" : { "32" : "ami-37c2f643", "64" : "ami-31c2f645" }  
    },  
    "Resources" : {  
        "myEC2Instance" : {  
            "Type" : "AWS::EC2::Instance",  
            "Properties" : {  
                "ImageId" : { "Fn::FindInMap" : [ "RegionMap", { "Ref" : "AWS::Region" }, "32" ]},  
                "InstanceType" : "m1.small"  
            }  
        }  
    }  
}
```



Functions

Fn::FindInMap

Fn::FindInMap: [*MapName*, *TopLevelKey*, *SecondLevelKey*]

Fn::GetAtt

Fn::GetAtt: [*LogicalNameOfResource*, *attributeName*]

Fn::GetAZs

Fn::ImportValue

Fn::ImportValue: *sharedValueToImport*

Fn::Join

Fn::Join: [*delimiter*, [*comma-delimited List of values*]]

Fn::Select

Fn::Select: [*index*, *ListofObjects*]

Fn::Split

Fn::Split: [*delimiter*, *source string*]

Fn::Sub

Fn::Transform

Ref

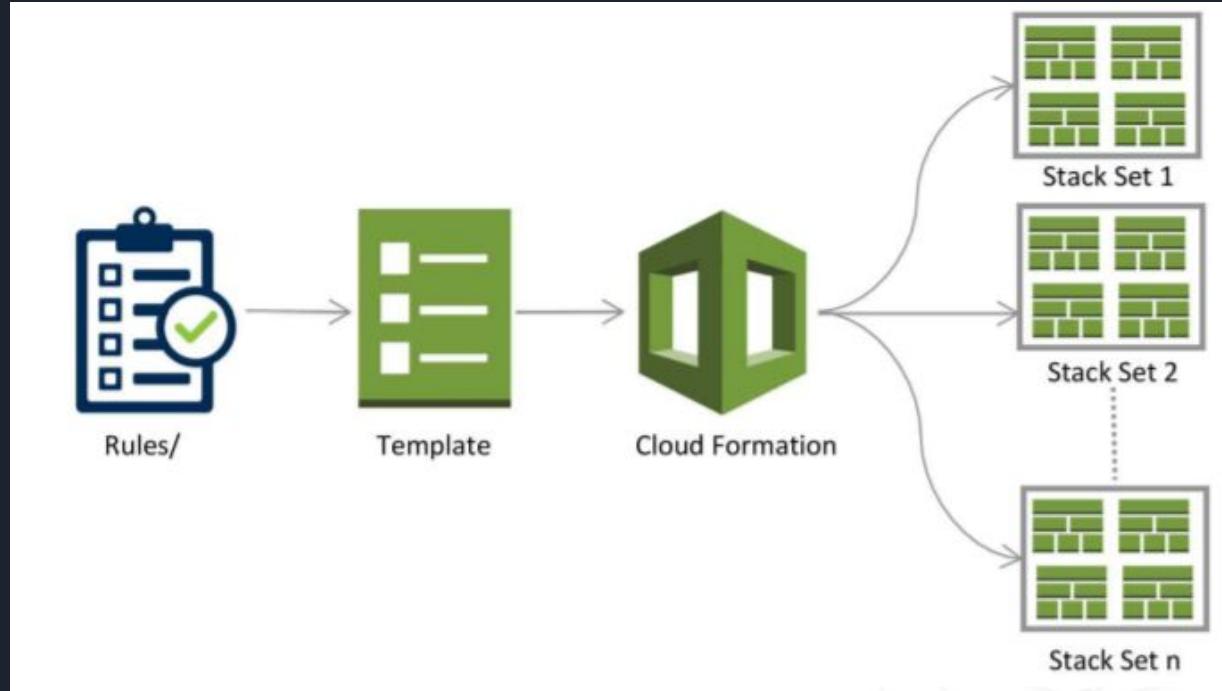
Ref: *LogicalName*

Fn::Sub:

- *String*
- *Var1Name*: *Var1Value*
- *Var2Name*: *Var2Value*

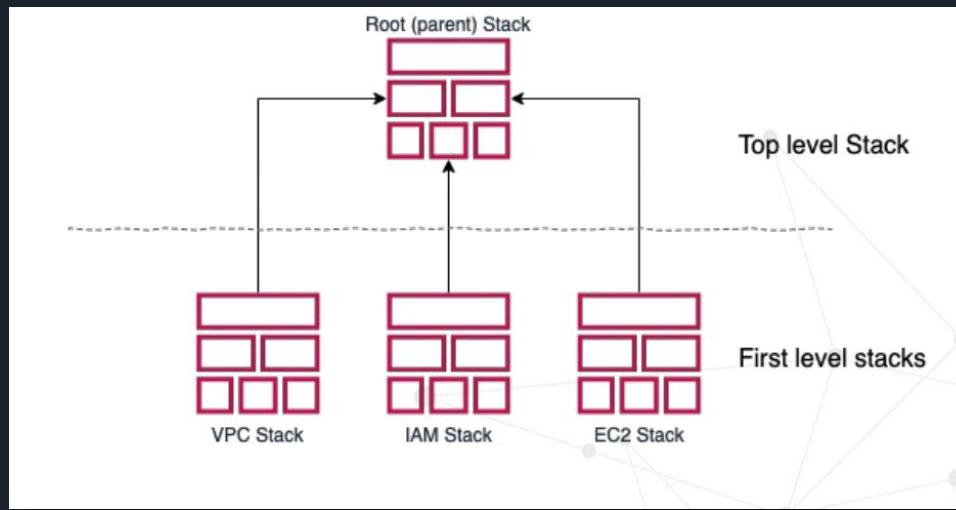
CloudFormation Stacks

- A collection of AWS resources is called a Stack and it can be managed in a single unit



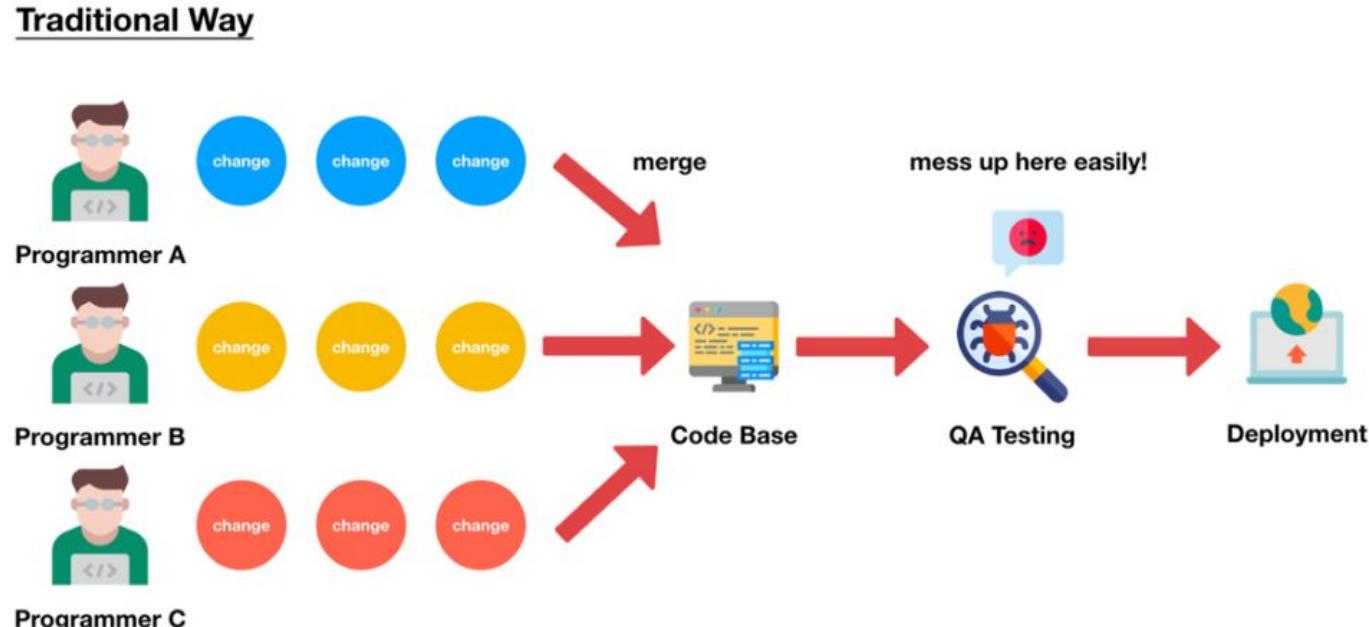
Nested Stacks

- Nested stacks are stacks created as part of other stacks. We create a nested stack within another stack by using the AWS::CloudFormation::Stack resource
- Nested stacks can themselves contain other nested stacks, resulting in a hierarchy of stacks, as in the diagram below. The root stack is the top-level stack to which all the nested stacks ultimately belong.



DevOps on AWS

Traditional Way



Challenges

Lack of Automated and Secured Workflow Management Systems in Business organization



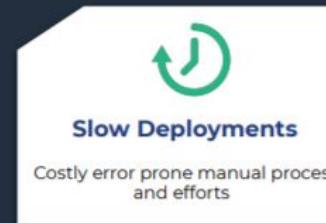
Building And Maintaining Servers

Time consuming and unproductive



No Environment Management

Differences in development and production environments



Slow Deployments

Costly error prone manual process and efforts



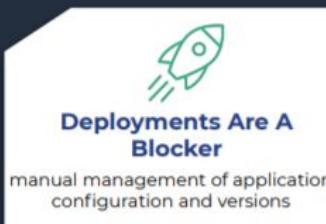
No Shared Ownership

Lack of feedback and proper metric leads



No Proper Configuration Management

Discrepancies in managing configurations



Deployments Are A Blocker

Manual management of application configuration and versions



Production Downtime

Due to lack of improper deployment instructions / checklist



Hacking & Security Threat

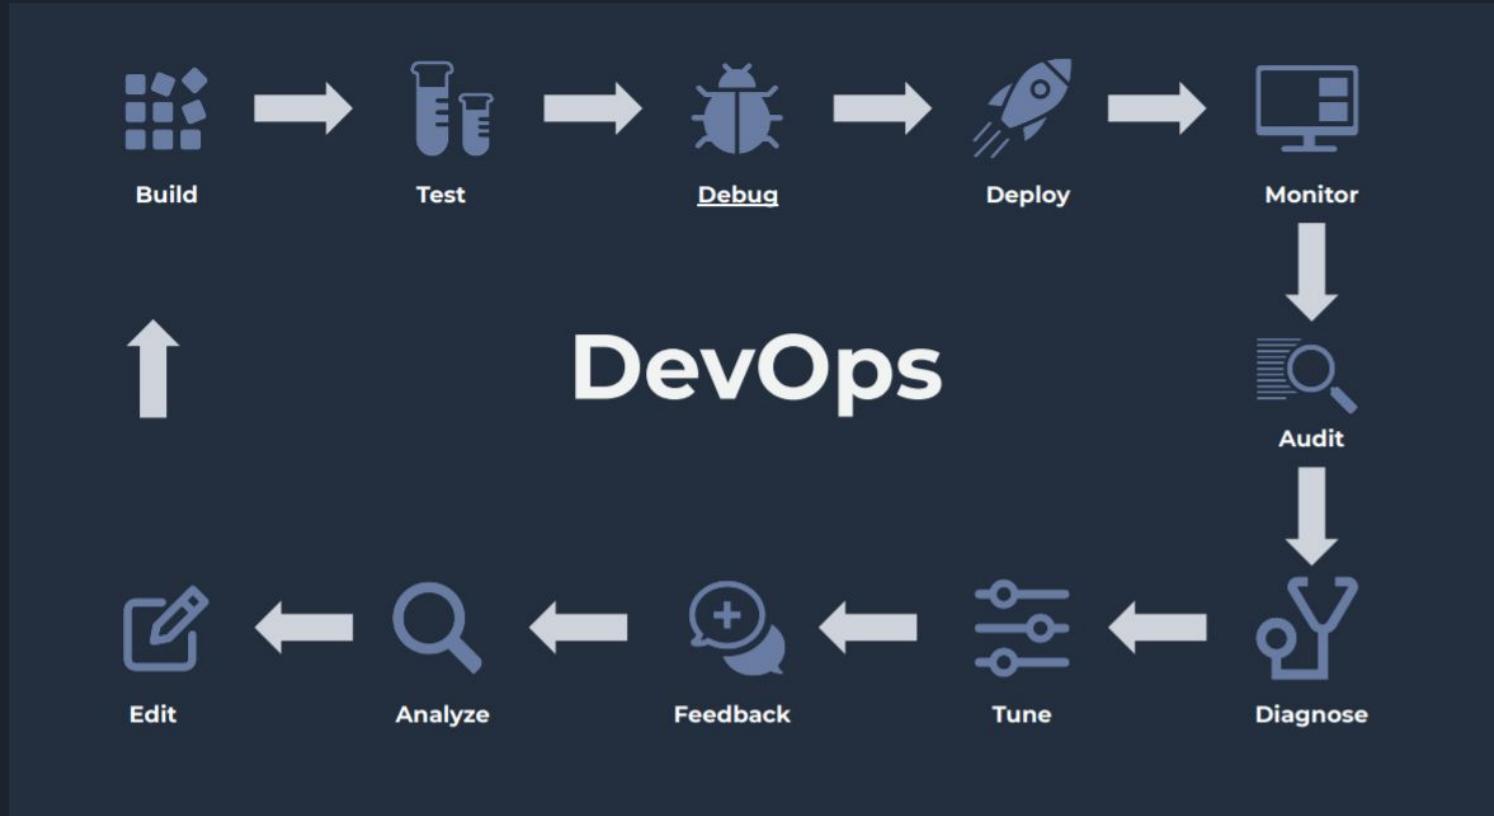
Fixing directly in production, instead of a proper hotfix process

Life Without Devops



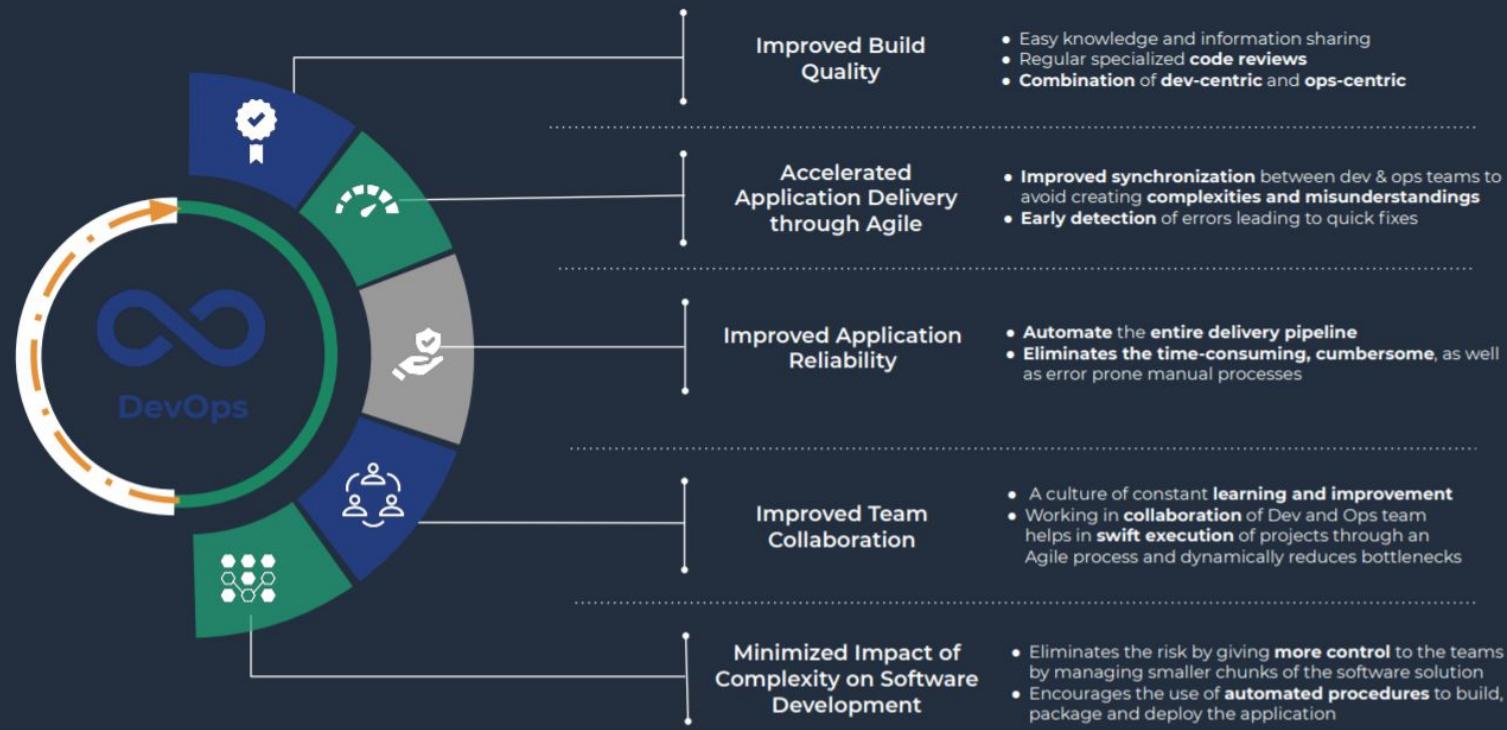
Meme Credits:
https://dev.to/presto412/hyperledger-fabric-transitioning-from-development-to-production-4dc_h

Devops Cycle



Devops Overview

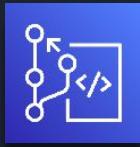
Business Impact by implementing DevOps



CICD Overview

- Continuous Integration
- Continuous Deployment
- Benefits of CICD
 - Small and frequent code changes
 - Faster Release Rate
 - Easier Rollbacks
 - Increase Team Transparency and Accountability

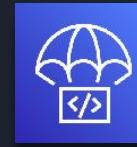
AWS CICD Tools



AWS
CodeCommit



AWS
CodeBuild



AWS
CodeDeploy



AWS
CodePipeline

AWS CodeCommit Overview

AWS CodeCommit Features



- Encryption
- Access Control using IAM
- Unlimited Repositories
- Easy Access and Integration
- Notifications and Custom Scripts

AWS CodeCommit Benefits



- Fully managed
- Secure
- High Availability and Durability
- Collaborate on code
- Faster development lifecycle
- Use your existing tools

CodeBuild Overview

- Fully Managed Build service(alternative to tools like jenkins)
- Serverless
- Secure
- Continuous scaling(no servers to manage)
- Pay only for the build time.
- Supports custom Docker images
- Build instructions defined in buildspec.yml file

buildspec.yml structure

```
version: 0.2
phases:
  install:
    commands:
      - echo Entered the install phase...
      - apt-get update -y
      - apt-get install -y maven
    finally:
      - echo This always runs even if the update or install command fails

  pre_build:
    commands:
      - echo Logging in to Amazon ECR...
      - aws ecr get-login-password --region $AWS_DEFAULT_REGION | docker login --username AWS --password-stdin
        $AWS_ACCOUNT_ID.dkr.ecr.$AWS_DEFAULT_REGION.amazonaws.com

  build:
    commands:
      - echo Build started on `date`
      - echo Building the Docker image...
      - docker build -t $IMAGE_REPO_NAME:$IMAGE_TAG .
      - docker tag $IMAGE_REPO_NAME:$IMAGE_TAG $AWS_ACCOUNT_ID.dkr.ecr.$AWS_DEFAULT_REGION.amazonaws.com/$IMAGE_REPO_NAME:$IMAGE_TAG
```

buildspecs.yaml structure

```
post_build:
  commands:
    - echo Build completed on `date`
    - echo Pushing the Docker image...
    - docker push $AWS_ACCOUNT_ID.dkr.ecr.$AWS_DEFAULT_REGION.amazonaws.com/$IMAGE_REPO_NAME:$IMAGE_TAG

artifacts:
  files:
    - target/messageUtil-1.0.jar

cache:
  paths:
    - '/root/.m2/**/*'
```

[link: buildspec.yml syntax](#)

CodeDeploy Overview

- Automates application deployment to EC2,Lambda,ECS
- Also supports deployment On-Premise servers
- Secure,Highly scalable,Fully managed
- Deploy instructions defined in appspec.yml file
- Automated Rollback

Note: CodeDeploy does not provision resources

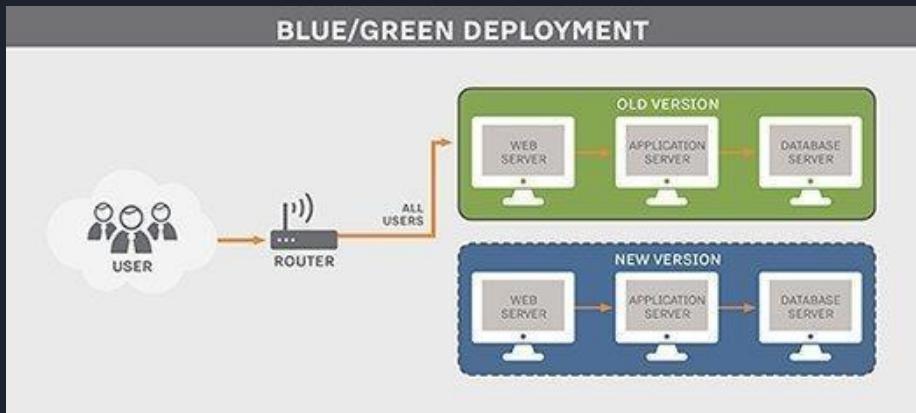
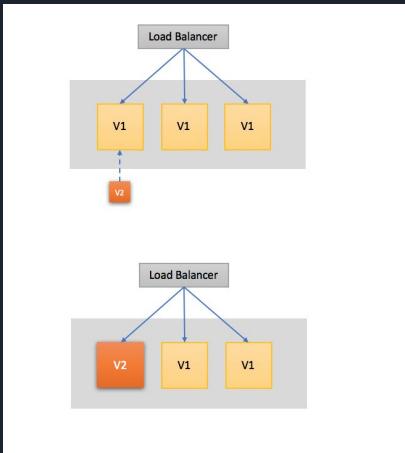
Working of CodeDeploy

Note: Install Codedeploy agent on EC2 or On Premise machines

- Code Deploy sends appspec.yml file to each machine
- Machines deploys the application based on instruction written in appspec.yml file.
- Code Deploy agent reports success/failure of deployment.

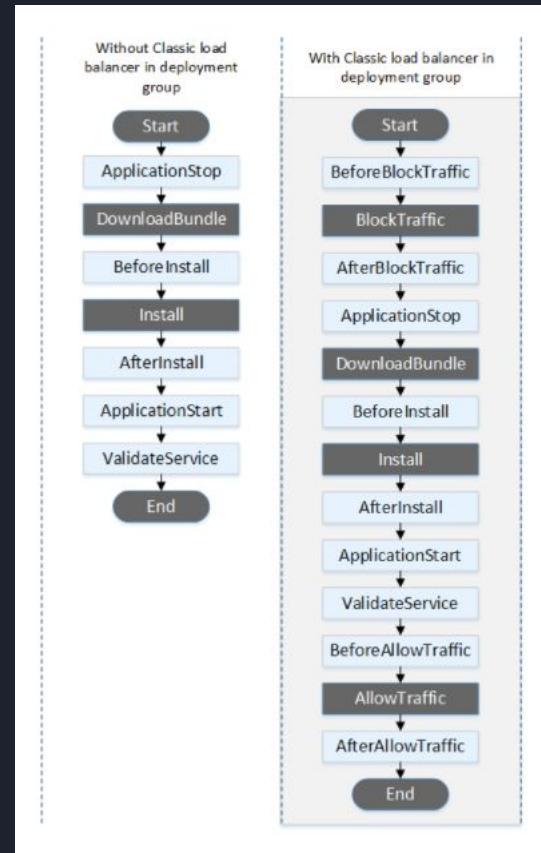
Deployment Type

- In place
- Blue/Green

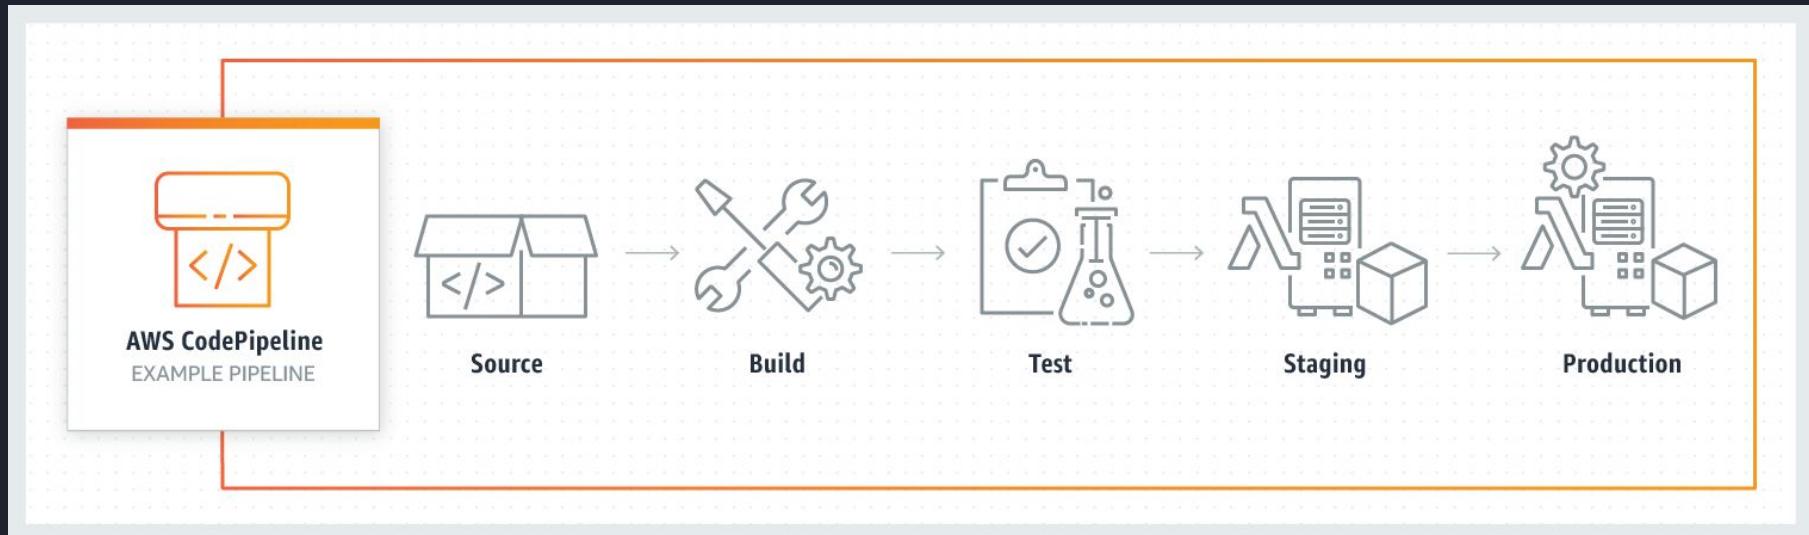


appspec.yml structure

```
version: 0.0
os: linux
files:
  - source: /
    destination: /var/www/html/WordPress
hooks:
  BeforeInstall:
    - location: scripts/install_dependencies.sh
      timeout: 300
      runas: root
  AfterInstall:
    - location: scripts/change_permissions.sh
      timeout: 300
      runas: root
  ApplicationStart:
    - location: scripts/start_server.sh
    - location: scripts/create_test_db.sh
      timeout: 300
      runas: root
  ApplicationStop:
    - location: scripts/stop_server.sh
      timeout: 300
      runas: root
```



CodePipeline Overview



<https://aws.amazon.com/codepipeline/>

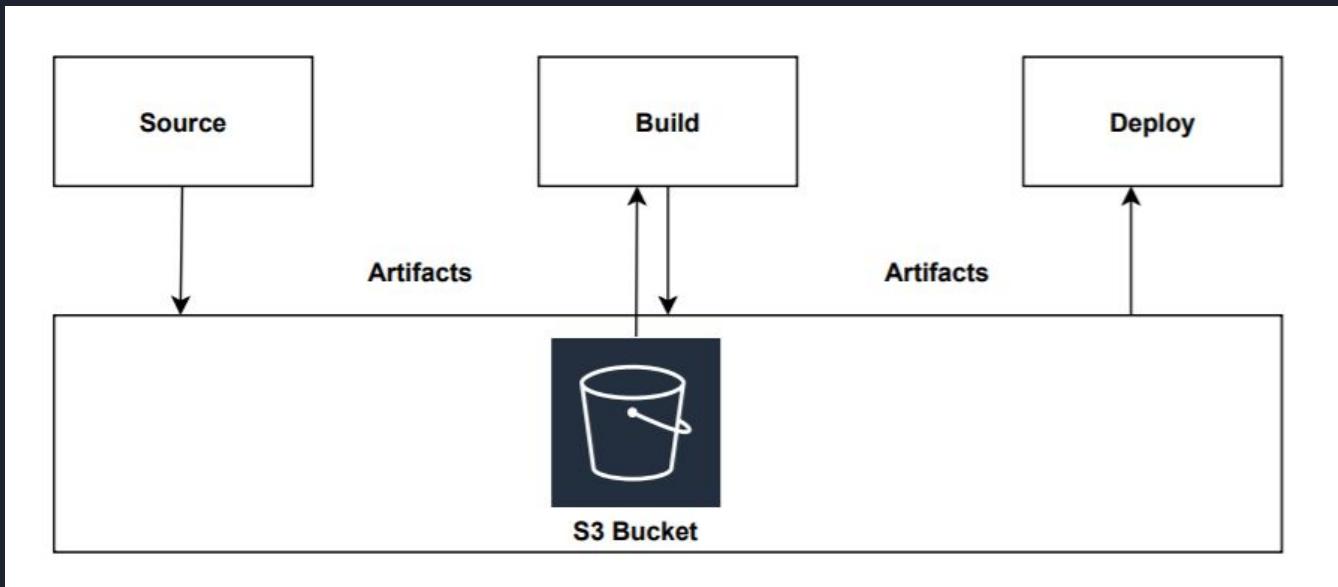
CodePipeline Overview

- Orchestrate entire CI/CD pipeline
- Visual workflow
- Multiple Stages
- Manual Approval
- Supports cross account and cross region deployment
- Pricing: 1 \$ per pipeline

<https://docs.aws.amazon.com/codepipeline/latest/userguide/welcome-introducing.html>

CodePipeline Working

- Each pipeline stage creates **artifacts**
- These artifacts are stored in s3 bucket.
- Each stage takes artifacts stored in s3 as input and dumps as output for next stage.



AWS CloudWatch

Cloudwatch Logs



Image Credits:
<https://www.pinterest.com/pin/743094007243620223/>

Cloudwatch Logs Overview

- Logging is essential to understand the behaviour of the application and to debug the issues which may impact business.
- Cloudwatch logs is used to collect log files
 - Logs can be used to review any event within a system, including failures, state transformations.

```
2020-11-11 13:52:15 INFO app - Loading configurations..
2020-11-11 13:52:18 INFO app - *** InstanceID API_V2_I02
2020-11-11 13:52:18 INFO app - *** BaseURL http://10.244.2.168:3000
2020-11-11 13:52:19 INFO app - *** LogLevel 05 (DEBUG)
2020-11-11 13:52:12 DEBUG App:22 - Initializing Swagger UI..
2020-11-11 13:52:14 DEBUG App:28 - Generating schemata..
2020-11-11 13:52:14 DEBUG App:30 - Initializing REST services..
2020-11-11 13:52:15 DEBUG App:32 - Generating documentation..
2020-11-11 13:52:18 DEBUG App:36 - Loading adapters..
2020-11-11 13:52:21 DEBUG Adapters:10 - Loading adapter docs/v1 from path: .....
2020-11-11 13:52:22 DEBUG Adapters:16 - Loading adapter mongo/v1 from path: .....
2020-11-11 13:52:26 DEBUG Docs:38 - document-store adapter initialized
2020-11-11 13:52:27 DEBUG Mongo:38 - mongo adapter initialized
2020-11-11 13:52:28 DEBUG Adapters:38 - Adapter: docs/v1 (Docs) v1.0.0 initialized
2020-11-11 13:52:28 DEBUG Adapters:38 - Adapter: mongo/v1 (Mongo) v1.0.0 initialized
2020-11-11 13:52:31 INFO app - Started listening...
2020-11-11 13:52:27 ERROR Oracle:22 - Failed to write to DB. Error: ORA-01017
2020-11-11 13:52:27 ERROR Users:24 - Failed to create JSON (E: missing user id)
```

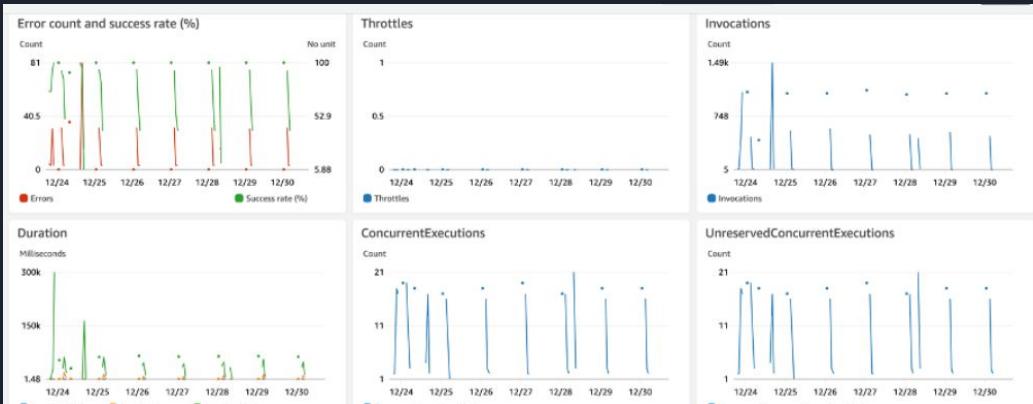
Cloudwatch Logs Overview

- In the production environment, we can't debug issues without proper log files as they become the only source of information
- Each service has one log group
- Logs can be queried
- Application specific Alarms can be created
- Cloudwatch agents for application logs
- Logs can be encrypted using KMS

Cloudwatch Metrics Overview

- Time-ordered set of data points published to CloudWatch
- Monitoring Types:

- Basic Monitoring
 - Collects data every 5 minutes
- Detailed Monitoring
 - Collects data every 1 minutes

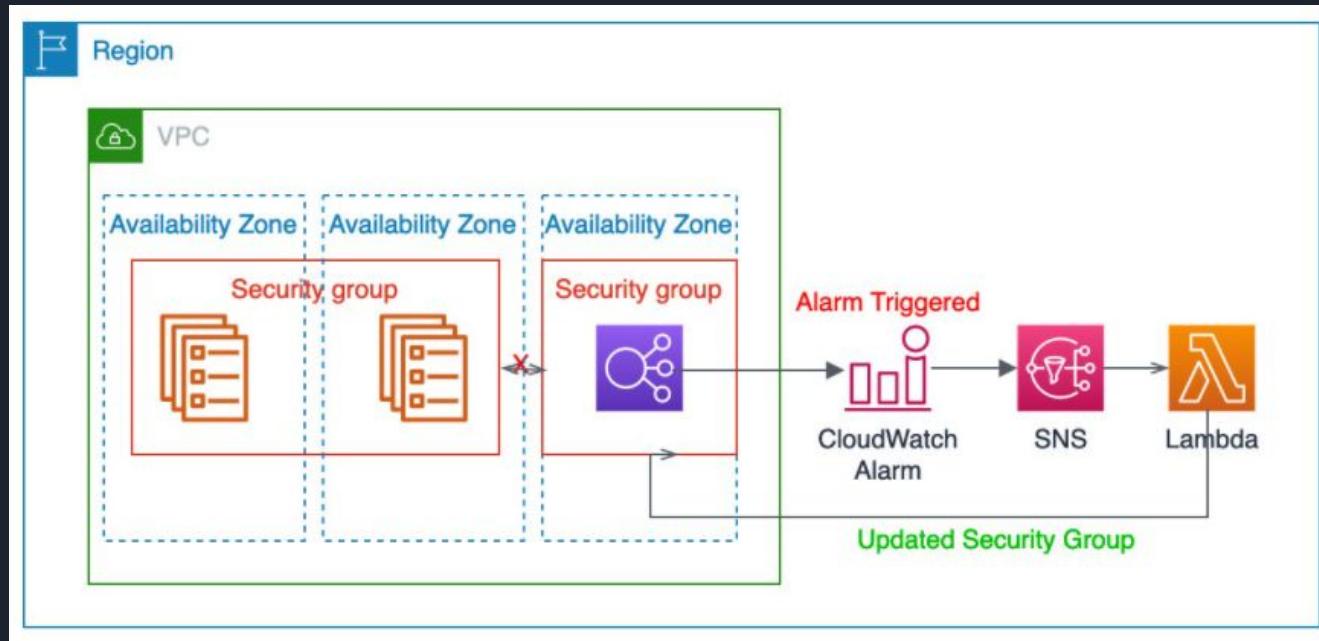


Cloudwatch Alarm Overview

- Alarm State:
 - In alarm
 - Insufficient data
 - OK
- Actions
 - SNS
 - EC2 Action
 - Auto scaling action

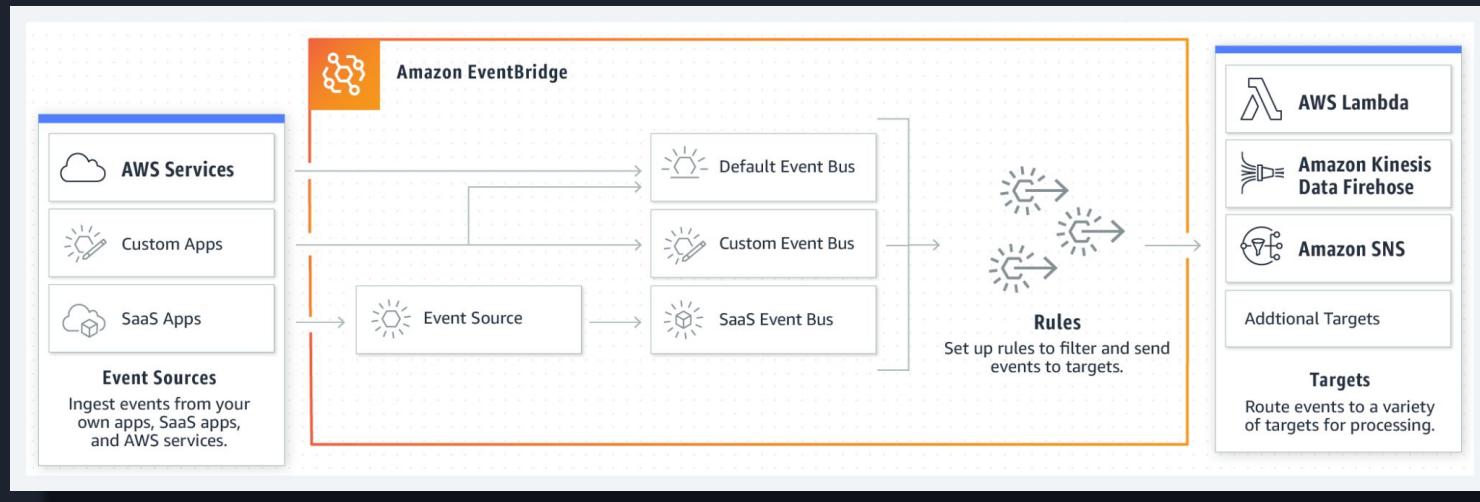


Automation using cloudwatch Alarm



EventBridge Overview

- Monitors the real time changes in AWS Services
- Using Rules, AWS services can act automatically when changes occur in other AWS Services
- Pay for what you use



AWS Cloudtrail

AWS CloudTrail Overview

```
"userIdentity": {  
    "accessKeyId": "AKEXAMPLE123EJVA",  
    "accountId": "123456789012",  
    "arn": "arn:aws:iam::123456789012:user/Bob",  
    "principalId": "AIEEXAMPLE987ZKLALD3HS",  
    "type": "IAMUser",  
    "userName": "Bob"  
}
```

What was the API call?

What resources were acted up on?

- API call and the service the API call belongs to

"eventName": "RunInstances"

"eventSource": "EC2"

Where was the API call made from?

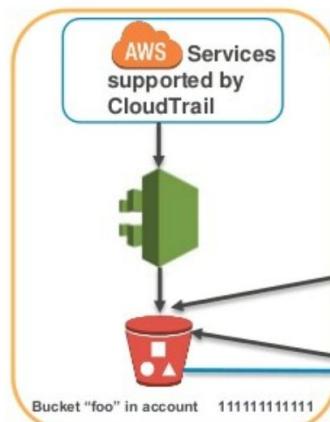
- Apparent IP address of the requester making the API call.
- Records the apparent IP address of the requester when making API calls from AWS Management Console.
- AWS region to which the API call was made. Global services (Examples: IAM/STS) will be recorded as us-east-1.

"sourceIPAddress": "54.234.127.135",
"awsRegion": "us-east-1",

AWS CloudTrail Overview

Aggregate log files from multiple accounts into one S3 bucket

1. Turn on CloudTrail for 111111111111



2. Update "foo" bucket policy

```
"arn:aws:s3:::foo/KBJInc/AWSLogs/222222222222/*",
"arn:aws:s3:::foo/KBJInc/AWSLogs/333333333333/*"
```

3. Turn on CloudTrail for 222222222222



4. Turn on CloudTrail for 333333333333



Thank you!

